

Designing a JPEG DCT in C/C++ using Catapult Synthesis

June 2017

**© 2015-17 Mentor Graphics Corporation
All rights reserved.**

This document contains information that is proprietary to Mentor Graphics Corporation. The original recipient of this document may duplicate this document in whole or in part for internal business purposes only, provided that this entire notice appears in all copies. In duplicating any part of this document, the recipient agrees to make every reasonable effort to prevent the unauthorized use and distribution of the proprietary information.

This document is for information and instruction purposes. Mentor Graphics reserves the right to make changes in specifications and other information contained in this publication without prior notice, and the reader should, in all cases, consult Mentor Graphics to determine whether any changes have been made.

The terms and conditions governing the sale and licensing of Mentor Graphics products are set forth in written agreements between Mentor Graphics and its customers. No representation or other affirmation of fact contained in this publication shall be deemed to be a warranty or give rise to any liability of Mentor Graphics whatsoever.

MENTOR GRAPHICS MAKES NO WARRANTY OF ANY KIND WITH REGARD TO THIS MATERIAL INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE.

MENTOR GRAPHICS SHALL NOT BE LIABLE FOR ANY INCIDENTAL, INDIRECT, SPECIAL, OR CONSEQUENTIAL DAMAGES WHATSOEVER (INCLUDING BUT NOT LIMITED TO LOST PROFITS) ARISING OUT OF OR RELATED TO THIS PUBLICATION OR THE INFORMATION CONTAINED IN IT, EVEN IF MENTOR GRAPHICS HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

U.S. GOVERNMENT LICENSE RIGHTS: The software and documentation were developed entirely at private expense and are commercial computer software and commercial computer software documentation within the meaning of the applicable acquisition regulations. Accordingly, pursuant to FAR 48 CFR 12.212 and DFARS 48 CFR 227.7202, use, duplication and disclosure by or for the U.S. Government or a U.S. Government subcontractor is subject solely to the terms and conditions set forth in the license agreement provided with the software, except for provisions which are contrary to applicable mandatory federal laws.

TRADEMARKS: The trademarks, logos and service marks ("Marks") used herein are the property of Mentor Graphics Corporation or other parties. No one is permitted to use these Marks without the prior written consent of Mentor Graphics or the owner of the Mark, as applicable. The use herein of a third- party Mark is not an attempt to indicate Mentor Graphics as a source of a product, but is intended to indicate a product from, or associated with, a particular third party. A current list of Mentor Graphics' trademarks may be viewed at: www.mentor.com/trademarks.

The registered trademark Linux® is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a world-wide basis.

End-User License Agreement: You can print a copy of the End-User License Agreement from: www.mentor.com/eula.

Mentor Graphics Corporation
8005 S.W. Boeckman Road, Wilsonville, Oregon 97070-7777.
Telephone: 503.685.7000
Toll-Free Telephone: 800.592.2210
Website: www.mentor.com
SupportNet: supportnet.mentor.com/

Improving Performance and Area via Design Partitioning

The previous Toolkits optimized the DCT for both performance and area using a combination of HLS constraints as well as architectural coding changes. This allowed us to achieve a throughput of 128 cycles and a large improvement in area via quantization using bit-accurate data types and architectural code changes (Folded filter). The use of a singleport RAM as well as the coding of sequential loops limits the performance on the order of 128 cycles. We saw from simulation that the ROW0 set of loops is idle when the COL1 set of loops is running and vice versa. The design loops must be partitioned into separate processes in order to achieve the final performance goal of 64 cycles. Catapult supports two type of design partitioning:

- Loosely coupled
 - Catapult “hierarchy” allows C++ functions to be synthesized as separate concurrently running processes with built-in synchronization.
- Tightly coupled
 - Catapult “CCORE” allows C++ functions to be synthesized as operators which can then be scheduled, shared, etc.

Compile the Design

- CD into the DCT_Hierarchy directory and launch Catapult by typing “catapult” at the command prompt (This step can be skipped if running the toolkit interactively).
- Go to File > Run Script and run the directives1.tcl file. If running interactively select the “Compile the design” script and click on “Launch Project” in the toolkit window.

Analyze the Code Changes for Catapult Hierarchy

- Open the “dct.cpp” file.
- Scroll through the design and look at how the code has been structured for Catapult “Hierarchy”.
 - ac_channel on interface declared as a reference “&”.
 - Arrays must be passed through channels via structs.
 - The top-level design only instantiates functions mapped to hierarchy and static interconnect channels

```
struct memStruct{ //Container for array passed through channel
    ac_int<16> data[8][8];
};
void dct(ac_channel<ac_int<8> > &input, ac_channel<memStruct > &output)
{
    static ac_channel<memStruct > mem; //Static interconnect channel

    //Top-level design only instantiates hierarchical blocks
    dct_h(input, mem);
    dct_v(mem, output);
}
```

- The ROW0 and COL1 set of loops have been moved into functions “dct_h” and “dct_v”
- The “dct_h” and “dct_v” functions have the hierarchy pragma “#pragma design” indicating that this function is mapped to hierarchy (This can also be done as a constraint).
- “dct_h” and “dct_v” both use the required coding style for memories, where the array mapped to memory must be in a struct (memStruct).
- “dct_h” and “dct_v” both define local structs of type “memStruct” to perform the memory reads/writes.

```
#pragma design
void dct_h(ac_channel<ac_int<8> > &input, ac_channel<memStruct > &mem)
    memStruct local_mem; //Local struct for temporary array operations
    ac_int<8> buf0[8]; //Local storage for one row of input
#pragma design
void dct_v(ac_channel<memStruct > &mem, ac_channel<memStruct > &output)
    memStruct local_mem; //Local struct for temporary array operations
    memStruct local_output;
```

- “dct_h” and “dct_v” both use the recommended coding style for arrays mapped to memories in hierarchy. Using the recommended style guarantees that the local structs will be optimized away leaving only memory interfaces.

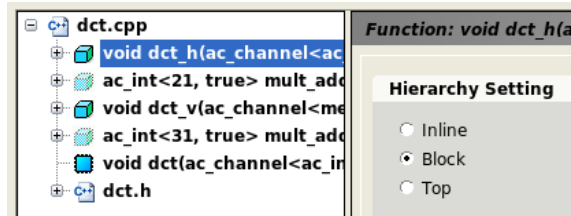
```
void dct_v(ac_channel<memStruct > &mem, ac_channel<memStruct > &output)
    memStruct local_mem; //Local struct for temporary array operations
    memStruct local_output;

    ac_int<16> buf1[8]; //Local storage for one row of input
    ac_int<31> acc1;


    local_mem = mem.read();
    COL1: for (int j=0; j < 8; ++j ) {
        COPY_ROW: for (int p=0; p<8; p++) //Copy one row of mem into local
            buf1[p] = local_mem.data[j][p];
        ROW1: for (int i=0 ; i < 8; ++i ){
            acc1 = mult_add<ac_int<16>, ac_int<17>, ac_int<31> > (buf1,i);
            local_output.data[i][j] = acc1 >> 15 ;
        }
    }
    output.write(local_output);
```

Constrain and Synthesize the Design


- Click on “Hierarchy” in the Task Bar.
- Note that the “dct_h” and “dct_v” Hierarchy Setting is set to “Block”




- Click on Mapping
- Note that “dct_h” and “dct_v” are shown as Instance Hierarchy. Their Design Type is set to “DESIGN” which indicates that they will be synthesized as separate concurrent processes.



- Go to Architectural Constraints and expand the Interface and Interconnect folders. You can see that the “output” channel and “mem” interconnect channel have been mapped to singleport RAM.



- Click on “dct_h” and unroll all the loops under the “ROW0” loop. Then pipeline the “main” loops with II=8.
- Repeat previous step for “dct_v”.



- Generate RTL
- Look at the Table View, the throughput is reported as 64 cycles.
- Run SCVerify and verify the throughput equals 64 cycles. Observe that “input” is read continuously and “output” is written continuously.

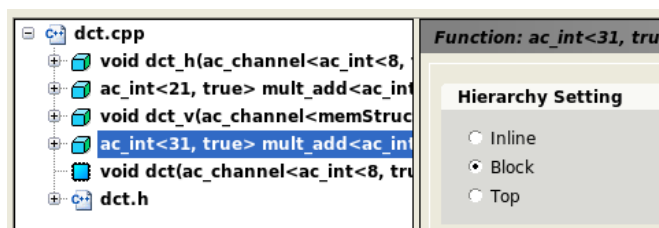
Optimize Area and Runtime using CCOREs

You may have noticed that the Catapult runtime slowed down as we added more parallelism into the single block design. Catapult's runtime is proportional to the number of operations in the design. Fully unrolling the inner loops and pipelining with $II=8$ creates lots of parallel operations, some of which Catapult can re-share because the $II=8$. However, doing this creates an “irregular” design, and Catapult will not be efficient when sharing resources (Note: you can usually see this irregularity in the Gantt chart). Catapult shares resources based on area cost. The bigger the resource the more likely it will get shared because Catapult has to weigh the area cost of adding input multiplexers to share a resource against the resource area itself. Unrolling, in this example, creates lots of smaller resources and Catapult is less efficient in sharing. Analysis of the C++ and constraints reveals that the fundamental computational unit for both “dct_h” and “dct_v” is the “mult_add” function. Although COL0 is unrolled 8 times, creating 8 copies of “mult_add”, “main” is pipelined with $II=8$, which should allow “mult_add” to be re-shared down to a single copy. In order for this to happen we need a way to tell Catapult to preserve “mult_add” as a single object so the entire object is easy to share. This can be accomplished using the Catapult CCORE flow.

```
void dct_h(ac_channel<ac_int<8> > &input, ac_channel<memStruct > &mem)
{
    memStruct local_mem; //Local struct for temporary array operations
    ac_int<8> buf0[8]; //Local storage for one row of input


    ac_int<21> acc0;
    ROW0: for (int i=0; i < 8; ++i) {
        COPY_ROW0: for (int p=0; p<8; p++) //Copy one row of input into local
            buf0[p] = input.read();
        COL0: for (int j=0; j < 8; ++j) {
            acc0 = mult_add<ac_int<8>,ac_int<9>,ac_int<21> >(buf0[j]);
            local_mem.data[j][i] = acc0 >> 5;
        }
    }
    mem.write(local_mem);
}
```

- Click on “Hierarchy” in the Task Bar.
- Set the Hierarchy Setting to “Block” for each instance of the “mult_add” function.



- Click on “Mapping” in the Task Bar.

- Select each instance of “mult_add” and set the “Block Type” to CCORE and the “CCORE Type” to “Combinational”.



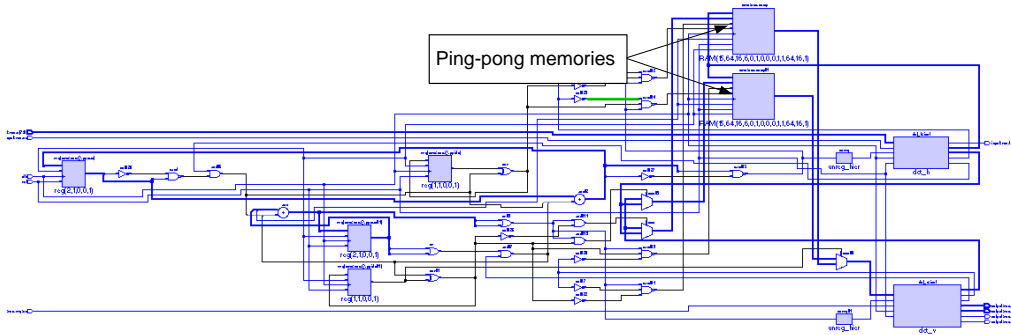
- Schedule the design.
- Open the Gantt chart and view it by component utilization.
- The schedule consists almost entirely of CCORE, memory, and IO operations. Note how regular the design “looks”.
- Generate RTL.
- Go to the Table View and look at how much area has been reduced.
- Set the Table View Report type to “Run Time”. Look at how much faster the design synthesized when using CCOREs. This is because most of the optimizations only have to deal with a single CCORE object as opposed to hundreds of individual operations.

Report: Run Time								
Solution /	Total	Analyze	Compile	Libraries	Assembly	Allocate	Schedule	Dpfsm
dct.v1 (extract)	239.13	3.13	1.33	0.35	0.39	2.22	226.73	4.98
dct.v2 (extract)	26.92	3.13	1.70	0.03	0.29	0.32	17.87	3.58

- Open the RTL report and look at the BOM. The entire design requires only 2 “mult_add” CCOREs, one for each “dct_h” and “dct_v”.

Bill Of Materials (Datapath)					
Component Name	Area	Score	Delay	Post Alloc	Post Assign
[Lib: cluster]					
mult_add_ac_int_16_true_ac_int_17_tru	10340.778	3.047		1	1
mult_add_ac_int_8_true_ac_int_9_true_	5377.327	3.143		1	1
[Lib: mac_inport]					

- Open the RTL schematic.
- Catapult has automatically created a ping-pong memory for the shared interconnect memory “mem”. This allows both blocks, dct_h and dct_v, to run simultaneously.

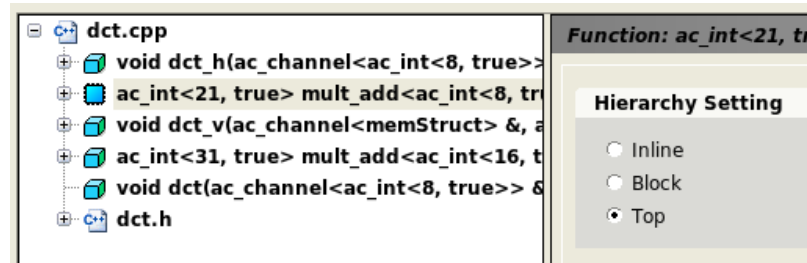


Bottom-up Design Flow


Using Catapult Hierarchy and CCOREs we were able to greatly improve design performance, area, and runtime. However, the design was being synthesized in a top-down fashion, which means that any C++ code change would force a re-synthesis of the entire design. This is bad for both runtime as well as impacting the back-end flows. Because of this Catapult allows designs to be synthesized in a bottom-up fashion at the DESIGN or CCORE level.

Synthesize the CCOREs

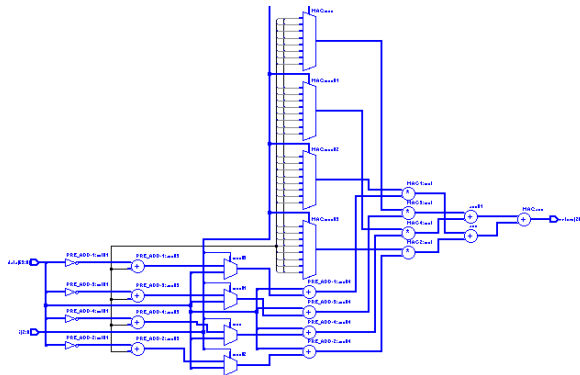
- Click on "Hierarchy in the Task Bar.
- Set the "ac_int<21,true> mult_add" function to be the top-level.



- GO to "Mapping" and set the Design Type to CCORE and the CCORE Type to Combinational.



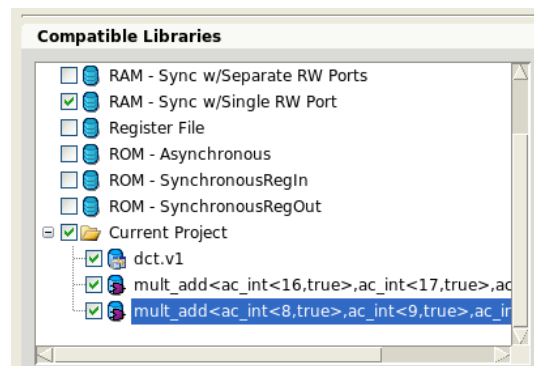
- Generate RTL.
- Open the RTL Schematic. Note the regularity of the mult_add function hardware.




- Go to mapping and repeat the previous steps on the "ac_int<31,true>mult_add" function.

Synthesize the Design Blocks

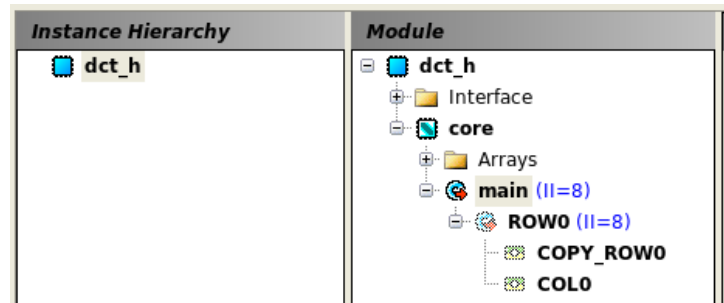
- Go to "Hierarchy" and set the "dct_h" function as the top-level design.
- Go to "Libraries" and select the two "mult_add" library components.




- Go to "Mapping", click on the "mult_add" instance, and set to Block Type to the pre-synthesized mult_add library component.




- Go to Architectural Constraints, unroll the COPY_ROW0 and COL0 loops and pipeline "main" with II=8.



- Generate RTL.
- Go to “Mapping” and make “dct_v” the top-level design.
- Repeat the previous steps for “dct_h” on “dct_v”.
- Go to “Mapping” and make “dct” the top-level design.



- Go to “Libraries” and note that “dct_h” and “dct_v” library components have been added to the list of libraries.
- GO to “Mapping” and map the “dct_h” and “dct_v” instances to their respective library component.



- Generate RTL. Note that this step is very fast since Catapult only needs to synthesize the interconnect and simply stitches the RTL already created for “dct_h” and “dct_v”.

DONE Toolkit